

Provider-Dependent Reliability of Self-Reported Confidence under a Shared Black-Box MT Protocol

Umberto Altieri

Abstract—This paper studies a trustworthy-AI reliability question for black-box language-model services: can one shared confidence threshold be reused across providers under a common protocol? We evaluate English→German translation with eight models from OpenAI, Anthropic, and Google on 500 WMT17 sentences (4,000 translation–confidence pairs) at temperature 0. Quality is scored with sentence-level chrF, and low quality is operationally defined as the within-model bottom 20% tail. Relative to that target, expected calibration error (ECE) and mismatch@0.9 differ widely despite similar mean chrF (54.0–60.1): ECE ranges from 0.076 to 0.180, and mismatch@0.9 ranges from 0.2% to 18.2%. The released snapshot also shows uneven parse-warning burden (443/4000 rows; 11.1%), an implementation-level confound for cross-provider reliability comparisons. The central claim is therefore practical and scoped: under this shared protocol and operational target, self-reported confidence is provider-dependent and non-portable as a common probability scale, so a fixed threshold does not induce comparable routing/abstention operating points. Supplementary analyses indicate that provider-specific isotonic remapping can reduce ECE but does not restore portability. Limitations are explicit: operational (not semantic) target, no bundled COMET/xCOMET checkpoints, no completed semantic-audit labels, and no bundled non-baseline prompt-variant outputs.

I. INTRODUCTION

We study whether prompt-elicited self-reported confidence can be reused across providers as a common threshold interface for MT routing/abstention decisions. The analysis is deliberately operational: low quality is defined by reference-overlap tails, not human semantic labels. The question is practical rather than theoretical: does one threshold (e.g., 0.9) induce comparable operating points across providers under a shared protocol?

Across eight models, mean chrF is relatively clustered while confidence behavior is not. This disconnect creates portability risk: identical threshold policies can produce materially different high-confidence failure rates. We therefore treat confidence as a provider-specific decision signal unless calibrated per provider.

II. RELATED WORK

MT reliability work spans reference-based evaluation (e.g., BLEU/chrF), learned metrics, and QE without references [10, 11, 12, 5, 13]. This paper does not propose a new metric or QE model; it asks whether the generating model’s own

confidence is stable enough for cross-provider threshold transfer [1].

In LLM settings, confidence and calibration have become central for selective prediction and abstention [3]. Prior work shows miscalibration in modern models and motivates post-hoc correction methods [6, 7, 14]; black-box settings are particularly challenging because token-level internals are unavailable [15].

Recent work on verbalized/self-reported confidence indicates useful signal but also sensitivity to elicitation and imperfect uncertainty expression [16]. Our contribution is a constrained MT portability test: same task, same protocol, multiple providers, and operational chrF/BLEU targets.

III. METHODS

Setup. We evaluate WMT17 en→de with a fixed 500-sentence sample (seed 123) [2]. Eight models (OpenAI/Anthropic/Google) each produce one translation and one scalar confidence in $[0, 1]$ per sentence (4,000 outputs total). Canonical prompts and JSON schema are shared; provider adapters only handle transport.

Operational target. Sentence-level chrF is the primary quality signal [11]. Low quality is the within-model bottom 20% chrF tail. This is an operational overlap-based target, not semantic correctness calibration. A secondary robustness stage uses the strongest available alternative metric in the environment (BLEU fallback when COMET is unavailable) [10, 12].

Metrics. With 10 bins, ECE is computed against the binary non-low-quality label, following standard calibration usage of ECE in modern ML [6]:

$$\text{ECE} = \sum_{b=1}^B \frac{|S_b|}{N} |\text{acc}(S_b) - \text{conf}(S_b)|.$$

Mismatch at threshold τ is

$$\Pr(\text{low-quality} \wedge \text{confidence} > \tau),$$

reported mainly at $\tau = 0.9$ (mismatch@0.9), a strict policy-relevant regime.

Complexity proxy. We include a lightweight source-side surface-complexity proxy (length/depth-like/rarity/name features) only for secondary stratification; it is not a validated semantic-difficulty measure.

Implementation details, full prompts, and script-level provenance are in the repository.

IV. RESULTS

Table I shows the core result. Mean chrF is similar (54.0–60.1), but ECE (0.076–0.180) and mismatch@0.9 (0.2%–18.2%) vary widely. Thus a fixed threshold does not transfer cleanly across providers.

Provider-specific isotonic remapping can reduce ECE for several models on held-out data, but gains are uneven and not portable; one model worsens in held-out ECE. This follows classic post-hoc probability calibration practice where isotonic mapping is model-specific [9]. This supports mitigation as provider-specific, not universal portability restoration.

Prompt-variant infrastructure and semantic-audit scaffolding are available in the repository, but non-baseline prompt outputs and completed semantic labels are not bundled in this release.

V. DISCUSSION AND LIMITATIONS

The main finding is comparative and operational: under one shared protocol, self-reported confidence behaves as a provider-dependent score, so fixed thresholds induce different routing/abstention points. This does not imply semantic calibration.

Limitations remain explicit: (i) target is operational chrF/BLEU overlap, not semantic correctness; (ii) no completed human semantic-audit labels; (iii) no bundled non-baseline prompt-variant outputs; (iv) one language pair/dataset with 500 sentences; (v) parse-warning burden (443/4000) is an implementation-level confound.

Full prompts, mismatch examples, robustness exports, and artifact-level provenance are included in the accompanying repository.

VI. CONCLUSION

Across 500 WMT17 sentences and 4,000 outputs from eight LLMs, confidence portability fails under a shared black-box MT protocol: similar mean chrF coexists with large cross-provider gaps in ECE and mismatch@0.9. The central claim is practical and scoped to this operational target.

Provider-specific isotonic calibration may mitigate alignment within providers, but it does not restore a shared probability scale across providers. Deployment should therefore use provider-specific threshold tuning/calibration or stronger external evaluator layers rather than assuming universal threshold portability; stronger learned evaluators are a natural next layer [8, 4].

REFERENCES

- [1] John Blatz, Erin Fitzgerald, George Foster, Simona Gandrabur, Cyril Goutte, Alex Kulesza, Alberto Sanchis, and Nicola Ueffing. Confidence estimation for machine translation. In *COLING 2004: Proceedings of the 20th International Conference on Computational Linguistics*, pages 315–321, Geneva, Switzerland, 2004. COLING. URL <https://aclanthology.org/C04-1046/>.
- [2] Ondřej Bojar, Christian Buck, Rajen Chatterjee, Christian Federmann, Yvette Graham, Barry Haddow, Matthias Huck, Antonio Jimeno Yepes, Philipp Koehn, Julia Kreutzer, et al. Findings of the 2017 conference on machine translation (wmt17). In *Proceedings of the Second Conference on Machine Translation*, pages 169–214, Copenhagen, Denmark, 2017. Association for Computational Linguistics. URL <https://aclanthology.org/W17-4717/>.
- [3] Ran El-Yaniv and Yair Wiener. On the foundations of noise-free selective classification. *Journal of Machine Learning Research*, 11:1605–1641, 2010. URL <https://jmlr.org/papers/v11/el-yaniv10a.html>.
- [4] Markus Freitag, Ricardo Rei, Nitika Mathur, Chi-kiu Lo, Craig Stewart, Eleftherios Avramidis, Tom Kocmi, George Foster, Alon Lavie, and André F. T. Martins. Results of wmt22 metrics shared task: Stop using bleu – neural metrics are better and more robust. In *Proceedings of the Seventh Conference on Machine Translation (WMT)*, pages 46–68, 2022. URL <https://aclanthology.org/2022.wmt-1.2/>.
- [5] Nuno M. Guerreiro, Ricardo Rei, Daan van Stigt, Luisa Coheur, Pierre Colombo, and André F. T. Martins. xcomet: Transparent machine translation evaluation through fine-grained error detection. *Transactions of the Association for Computational Linguistics*, 12: 979–995, 2024. doi: 10.1162/tacl_a_00683. URL <https://aclanthology.org/2024.tacl-1.54/>.
- [6] Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, pages 1321–1330, 2017. URL <https://proceedings.mlr.press/v70/guo17a.html>.
- [7] Zhengbao Jiang, Jun Araki, Haibo Ding, and Graham Neubig. How can we know when language models know? on the calibration of language models for question answering. *Transactions of the Association for Computational Linguistics*, 9:962–977, 2021. URL <https://aclanthology.org/2021.tacl-1.57/>.
- [8] Tom Kocmi and Christian Federmann. Large language models are state-of-the-art evaluators of translation quality. In *Proceedings of the 24th Annual Conference of the European Association for Machine Translation*, pages 193–203, 2023. URL <https://aclanthology.org/2023.eamt-1.19/>.
- [9] Alexandru Niculescu-Mizil and Rich Caruana. Predicting good probabilities with supervised learning. In *Proceedings of the 22nd International Conference on Machine Learning*, pages 625–632. ACM, 2005. doi: 10.1145/1102351.1102430.
- [10] Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: a method for automatic evaluation of machine translation. In *Proceedings of the 40th Annual*

Model	Mean chrF	ECE (within-q20)	Mismatch@0.9	Median conf.
anthropic/claude-haiku-4-5-20251001	57.01	0.094	13.4%	0.92
anthropic/claude-opus-4-6	58.45	0.087	7.6%	0.93
anthropic/claude-sonnet-4-5-20250929	56.89	0.095	12.8%	0.95
gemini/gemini-2.5-flash	54.01	0.148	4.2%	0.73
gemini/gemini-2.5-pro	60.07	0.180	18.2%	1.00
openai/gpt-5-mini	57.09	0.152	12.8%	0.93
openai/gpt-5-nano	55.31	0.076	0.2%	0.73
openai/gpt-5.2	58.87	0.134	15.2%	0.93

TABLE I

MAIN SUMMARY TABLE FROM REGENERATED CODE-SIDE OUTPUTS (WITHIN-MODEL CHRf-Q20 TARGET).

Model	ECE before	ECE after	Mismatch@0.9 before	Mismatch@0.9 after
anthropic/claude-haiku-4-5-20251001	0.110	0.027	14.3%	0.3%
anthropic/claude-opus-4-6	0.084	0.114	9.0%	2.1%
anthropic/claude-sonnet-4-5-20250929	0.087	0.073	12.1%	1.0%
gemini/gemini-2.5-flash	0.152	0.055	4.3%	0.3%
gemini/gemini-2.5-pro	0.175	0.087	17.5%	0.0%
openai/gpt-5-mini	0.128	0.116	10.8%	0.0%
openai/gpt-5-nano	0.091	0.045	0.3%	0.3%
openai/gpt-5.2	0.155	0.047	16.3%	0.0%

TABLE II

HELD-OUT ISOTONIC CALIBRATION SUMMARY FROM REGENERATED CALIBRATION ARTIFACTS.

- Meeting of the Association for Computational Linguistics*, pages 311–318, Philadelphia, Pennsylvania, USA, 2002. Association for Computational Linguistics. URL <https://aclanthology.org/P02-1040/>.
- [11] Maja Popović. chrF: character n-gram f-score for automatic mt evaluation. In *Proceedings of the Tenth Workshop on Statistical Machine Translation*, pages 392–395, Lisbon, Portugal, 2015. Association for Computational Linguistics. URL <https://aclanthology.org/W15-3049/>.
- [12] Ricardo Rei, Craig Stewart, Ana C. Farinha, and Alon Lavie. Comet: A neural framework for mt evaluation. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 2685–2702, 2020. URL <https://aclanthology.org/2020.emnlp-main.213/>.
- [13] Ricardo Rei, Marcos Treviso, Nuno M. Guerreiro, Ana C. Farinha, Duarte Alves, Charles Corbière, Eleftherios Avramidis, and André F. T. Martins. Cometkiwi: Ist-unbabel 2022 submission for the quality estimation shared task. In *Proceedings of the Seventh Conference on Machine Translation (WMT)*, pages 634–645, 2022. URL <https://aclanthology.org/2022.wmt-1.60/>.
- [14] Katherine Tian, Eric Mitchell, Allan Zhou, Archit Sharma, Rafael Rafailov, Huaxiu Yao, Chelsea Finn, and Christopher D. Manning. Just ask for calibration: Strategies for eliciting calibrated confidence scores from language models fine-tuned with human feedback. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 12903–12917, 2023. URL <https://aclanthology.org/2023.emnlp-main.330/>.
- [15] Dennis Ulmer, Martin Gubri, Hwaran Lee, Sangdoon Yun, and Seong Oh. Calibrating large language models using their generations only. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15440–15459, Bangkok, Thailand, 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.824. URL <https://aclanthology.org/2024.acl-long.824/>.
- [16] Gal Yona, Roei Aharoni, and Mor Geva. Can large language models faithfully express their intrinsic uncertainty in words? In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 7752–7764, Miami, Florida, USA, 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.emnlp-main.443. URL <https://aclanthology.org/2024.emnlp-main.443/>.




Fig. 1. Reliability diagram (qualitative) under the operational chrF-based label.